

Efficient Programming on ARM

Shiao-Li Tsao

Outline

- ARM System Software
 - ARM Bootstrap code
 - ARM OS code
 - Bootloader case study based on Uboot
 - Embedded OS case study based on uC/OS-II
- ARM Application Software
 - C compiler options
 - C programming skills on ARM
 - Assemble programming skills on ARM

Something Before Program Development

Slides are based on Vijay Kumar B., “Embedded
Programming with the GNU Toolchain”

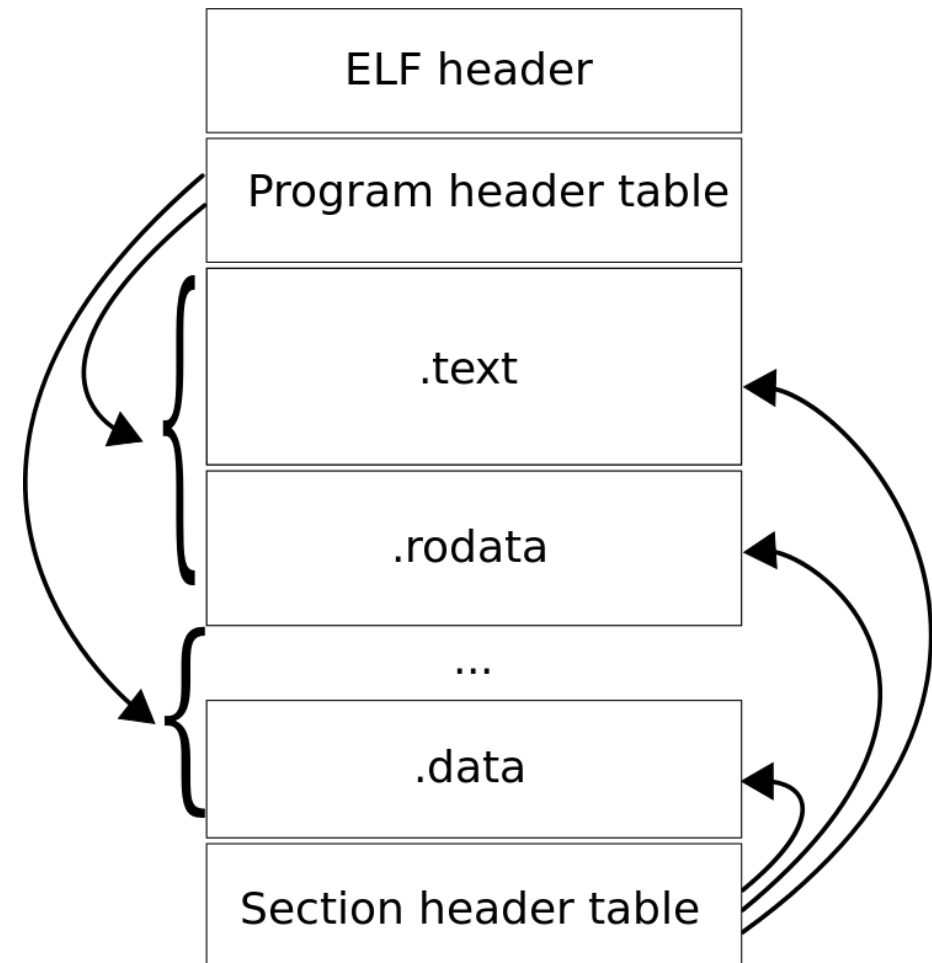

<http://www.bravegnu.org/gnu-eprog/>

Linking

```
1 // main.c
2 //
3 void sayHello();
4
5 char str[6] = {'h', 'e', 'l', 'l', 'o', 0};
6
7 int main(int argc, char *argv[])
8 {
9     sayHello();
10    return 0;
11 }
```

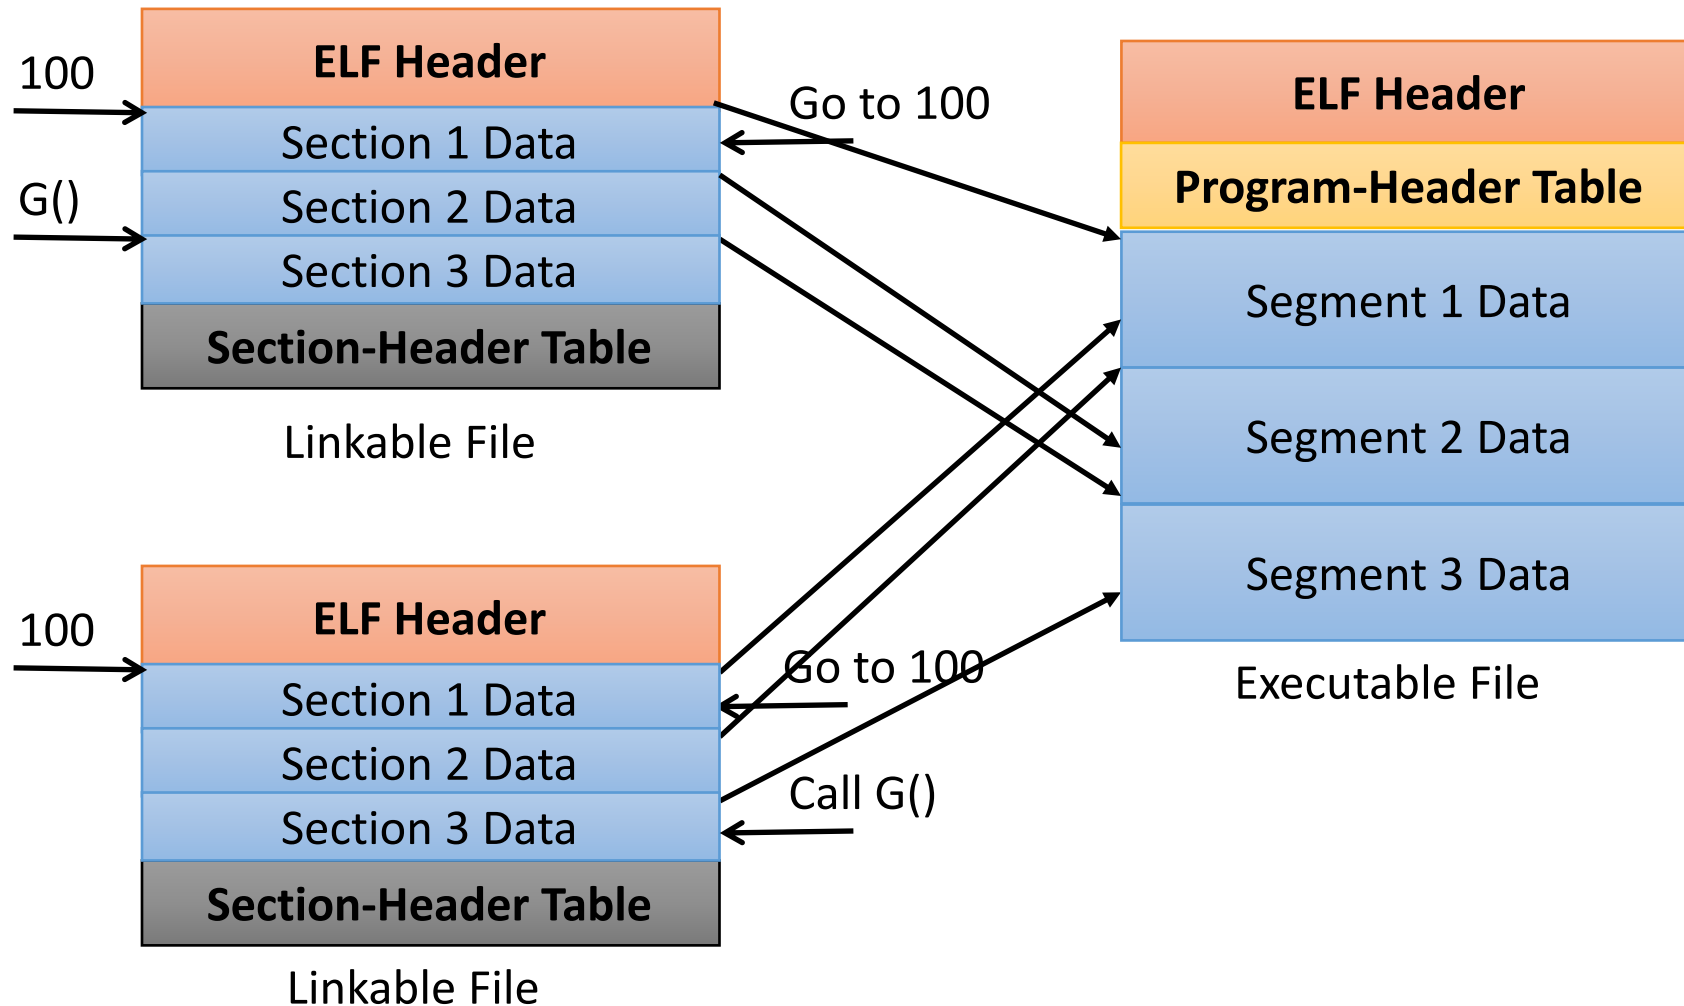
```
1 // hello.c
2 //
3 #include <stdio.h>
4
5 extern char str[6];
6
7 void sayHello() {
8     printf("%s\n", str);
9 }
10
11
```

Linking



By Surueña - Own work, CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=2922583>

Linking

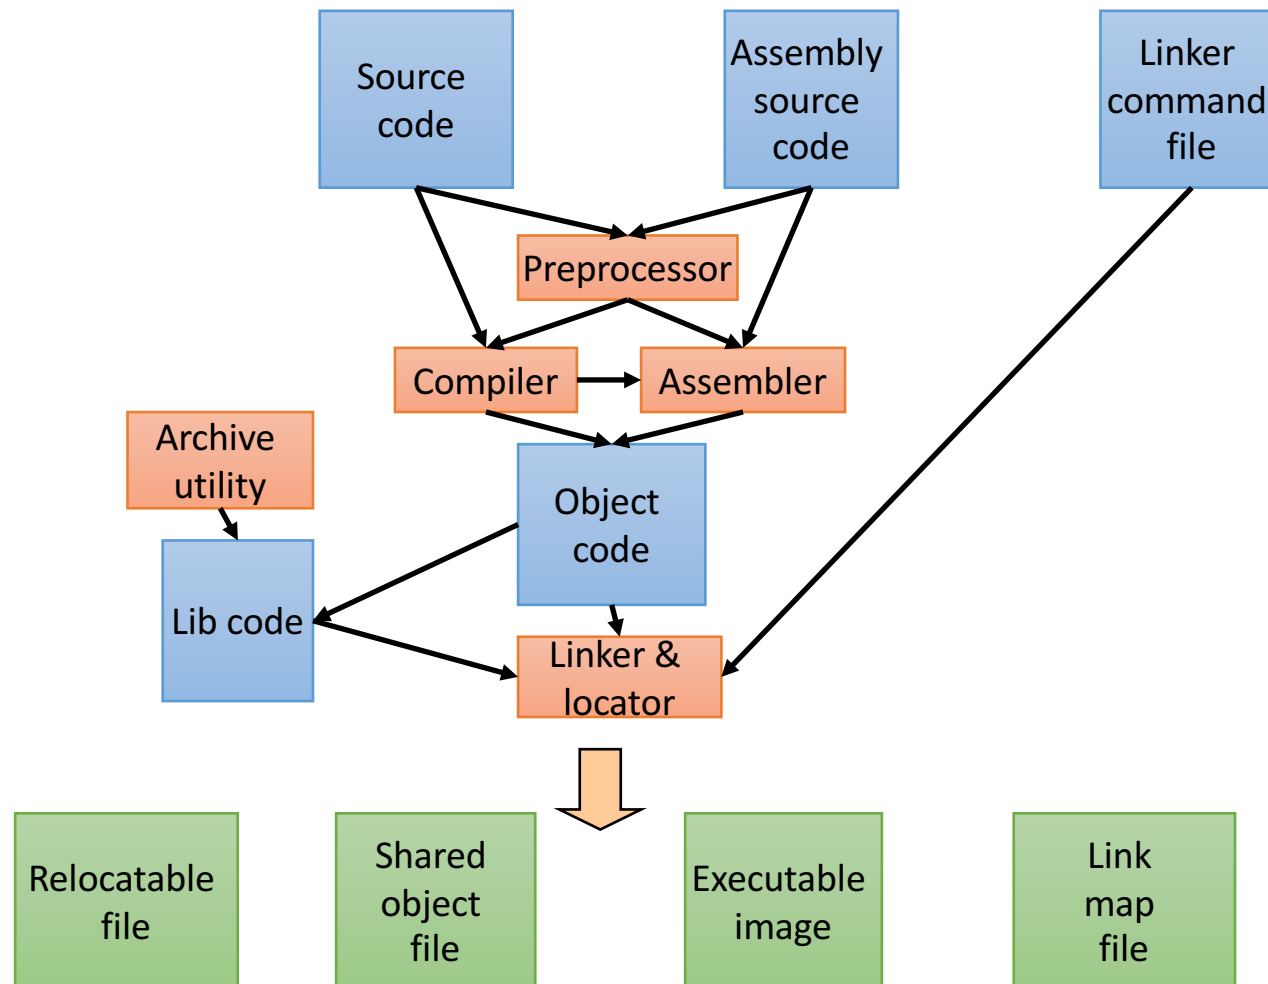


Linking

- Please review system programming and compiler if you are not familiar with below terms
 - Static linking
 - Dynamic linking
 - Relocations
 - Symbol table
 - Share library
 - Linking and loading

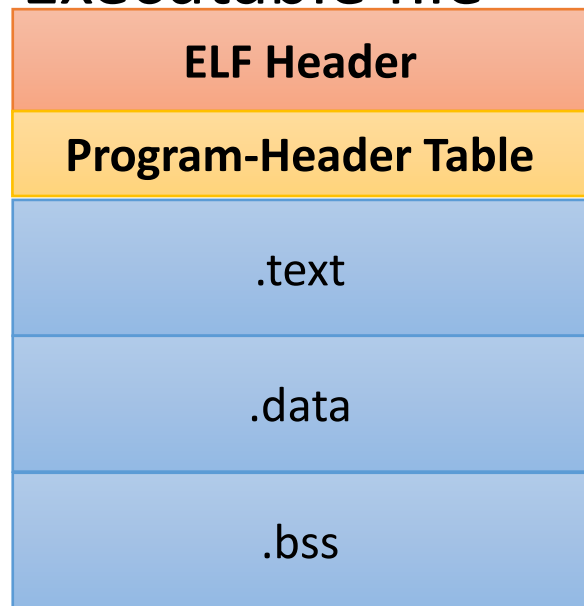
How to compile kernel codes

- Creating an executable image




Memory layout

Executable file



Executable File



Assembly, Machine code, Memory dump

Assembly ->

```
Dump of assembler code for function system_call:
0x000076d8 <+0>:    cmp     $0x47,%eax
0x000076db <+3>:    ja      0x76c8 <bad_sys_call>
0x000076dd <+5>:    push    %ds
0x000076de <+6>:    push    %es
0x000076df <+7>:    push    %fs
0x000076e1 <+9>:    push    %edx
=> 0x000076e2 <+10>:   push    %ecx
0x000076e3 <+11>:   push    %ebx
0x000076e4 <+12>:   mov     $0x10,%edx
0x000076e9 <+17>:   mov     %edx,%ds
0x000076eb <+19>:   mov     %edx,%es
0x000076ed <+21>:   mov     $0x17,%edx
0x000076f2 <+26>:   mov     %edx,%fs
0x000076f4 <+28>:   call    *0x1a020(,%eax,4)
0x000076fb <+35>:   push    %eax
0x000076fc <+36>:   mov     0x1b140,%eax
0x00007701 <+41>:   cmpl    $0x0,(%eax)
0x00007704 <+44>:   jne     0x76ce <reschedule>
0x00007706 <+46>:   cmpl    $0x0,0x4(%eax)
0x0000770a <+50>:   je      0x76ce <reschedule>
End of assembler dump.
```

Memory dump

|
v

0x76d8 <system_call>:	131	248	71	119	235	30	6	15	
0x76e0 <system_call+8>:	160	82	81	83	186	16	0	0	
0x76e8 <system_call+16>:		0	142	218	142	194	186	23	0
0x76f0 <system_call+24>:		0	0	142	226	255	20	133	32
0x76f8 <system_call+32>:		160	1	0	80	161	64	177	1
0x7700 <system_call+40>:		0	131	56	0	117	200	131	120
0x7708 <system_call+48>:		4	0						

Sections (Contd.)

```
.data
arr: .word 10, 20, 30, 40, 50
len: .word 5
.text
start: mov r1, #10
      mov r2, #20
      .data
result: .skip 4
      .text
      add r3, r2, r1
      sub r3, r2, r1
```

.data section

```
0000_0000 arr: .word 10, 20, 30, 40, 50
0000_0014 len: .word 5
0000_0018 result: .skip 4
```

.text section

```
0000_0000 start: mov r1, #10
0000_0004      mov r2, #20
0000_0008      add r3, r2, r1
0000_000C      sub r3, r2, r1
```

- Source – sections can be interleaved
- Bytes of a section – contiguous addresses

a.s (.text)

```
strcpy: ldrb r0, [r1], #1
        strb r0, [r2], #1
        cmp  r0, 0
        bne  strcpy
        mov  pc, lr
```

Assembler

```
0000_0000 strcpy: ldrb r0, [r1], #1
0000_0004         strb r0, [r2], #1
0000_0008         cmp  r0, 0
0000_000C         bne  strcpy
0000_0010         mov  pc, lr
```

b.s (.text)

```
strlen: ldrb r0, [r1], #1
        add  r2, #1
        cmp  r0, 0
        bne  strlen
        mov  pc, lr
```

Assembler

```
0000_0000 strlen: ldrb r0, [r1], #1
0000_0004         add  r2, #1
0000_0008         cmp  r0, 0
0000_000C         bne  strcpy
0000_0010         mov  pc, lr
```

```
0000_0000 strcpy: ldrb r0, [r1], #1
0000_0004         strb r0, [r2], #1
0000_0008         cmp r0, 0
0000_000C         bne strcpy
0000_0010         mov pc, lr
```

```
0000_0000 strlen: ldrb r0, [r1], #1
0000_0004         add r2, #1
0000_0008         cmp r0, 0
0000_000C         bne strlen
0000_0010         mov pc, lr
```

Merging .text sections from two files

```
0000_0000 strcpy: ldrb r0, [r1], #1
0000_0004         strb r0, [r2], #1
0000_0008         cmp r0, 0
0000_000C         bne strcpy
0000_0010         mov pc, lr
0000_0014 strlen: ldrb r0, [r1], #1
0000_0018         add r2, #1
0000_001C         cmp r0, 0
0000_0020         bne strlen
0000_0024         mov pc, lr
```

New address
after merge

Patched

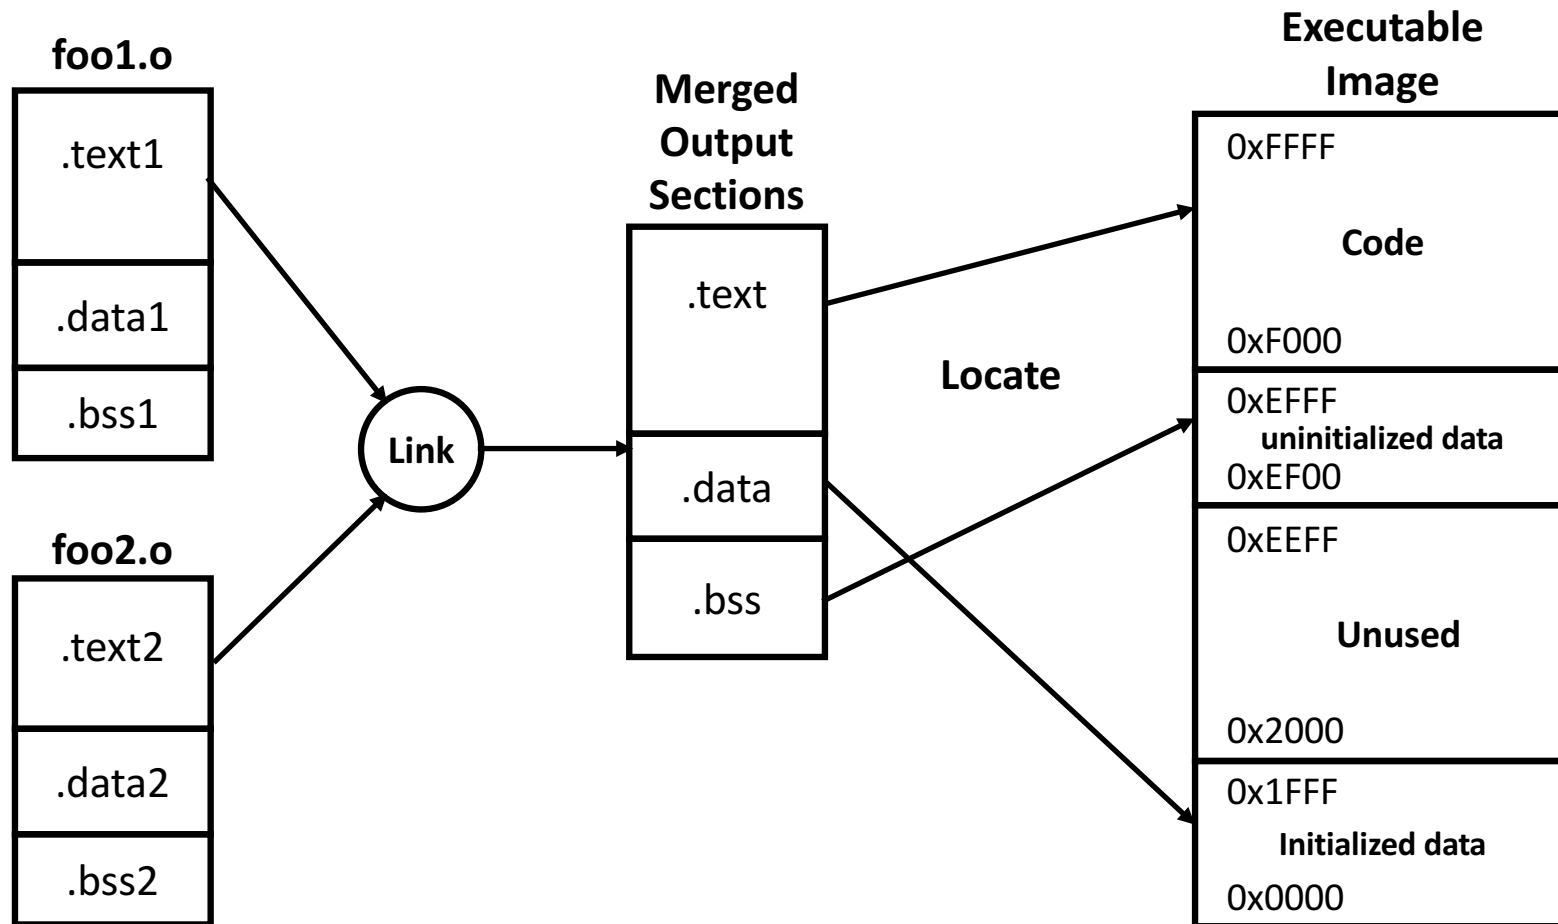
```
0000_0000 strcpy: ldrb r0, [r1], #1
0000_0004          strb r0, [r2], #1
0000_0008          cmp  r0, 0
0000_000C          bne  strcpy
0000_0010          mov  pc, lr
0000_0014 strlen: ldrb r0, [r1], #1
0000_0018          add  r2, #1
0000_001C          cmp  r0, 0
0000_0020          bne  strlen
0000_0024          mov  pc, lr
```

Placing .text section at 0x2000_0000

```
2000_0000 strcpy: ldrb r0, [r1], #1
2000_0004          strb r0, [r2], #1
2000_0008          cmp  r0, 0
2000_000C          bne  strcpy
2000_0010          mov  pc, lr
2000_0014 strlen: ldrb r0, [r1], #1
2000_0018          add  r2, #1
2000_001C          cmp  r0, 0
2000_0020          bne  strlen
2000_0024          mov  pc, lr
```

Patched


Linker and Locator Flows



Linker Script Revisited

```
MEMORY {  
    FLASH (rx) : ORIGIN = 0x00000000, LENGTH = 0x10000  
    SRAM (rwx) : ORIGIN = 0x20000000, LENGTH = 0x2000  
}
```

```
SECTIONS {  
    .text : {  
        * (.text);  
    } > FLASH  
  
    .data : {  
        * (.data);  
    } > SRAM  
}
```



Make and Makefile

- Make: utility to provide a convenient facility to build, install, and uninstall projects
- Makefile: script file for make to compile and link programs

Make and Makefile

```
target ... : prerequisites ...  
      recipe  
      ...  
      ...
```

```
edit : main.o kbd.o command.o display.o \  
      insert.o search.o files.o utils.o  
      cc -o edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
clean :  
      rm edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

Make and Makefile

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
clean :  
      rm edit $(objects)
```

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
  
main.o : defs.h  
kbd.o : defs.h command.h  
command.o : defs.h command.h  
display.o : defs.h buffer.h  
insert.o : defs.h buffer.h  
search.o : defs.h buffer.h  
files.o : defs.h buffer.h command.h  
utils.o : defs.h  
  
.PHONY : clean  
clean :  
      rm edit $(objects)
```

Make and Makefile

- Variables and settings
 - make config
 - make menuconfig
- Phony targets
 - make all
 - make clean
 - make depend
 - make install
 - make uninstall

The portability problem

- Hardware differences
- OS differences
- Compiler differences

```
your source files --> [autoscan*] --> [configure.scan] --> configure.ac
```

```
configure.ac --.  
              | .-----> autoconf* -----> configure  
[aclocal.m4] --+----+  
              | '-----> [autoheader*] --> [config.h.in]  
[acsite.m4] ---'
```

```
Makefile.in
```

```
[acinclude.m4] --.  
              |  
[local macros] --+---> aclocal* --> aclocal.m4  
              |  
configure.ac ----'
```

The portability problem

```
configure.ac --.  
               +--> automake* --> Makefile.in  
Makefile.am ---'  
  
               .-----> [config.cache]  
configure* -----+-----> config.log  
                  |  
[config.h.in] -.      v      .-> [config.h] -.  
               +--> config.status* -+      +--> make*  
Makefile.in ---'               '-> Makefile ---'
```

Minimal Bootloader

Minimal Bootloader Procedures

Step	Description
1	Take the Reset exception
2	Start initializing the hardware
3	Remap memory
4	Initialize communication hardware
5	Bootloader—copy payload and relinquish control

Step 1: Take the Reset Exception

```
AREA start, CODE, READONLY
ENTRY
```

```
sandstone_start
```

```
    B    sandstone_init1    ; reset vector
    B    ex_und              ; undefined vector
    B    ex_swi              ; swi vector
    B    ex_pabt             ; prefetch abort vector
    B    ex_dabt             ; data abort vector
    NOP                      ; not used...
    B    int_irq             ; irq vector
    B    int_fiq             ; fiq vector
```

```
ex_und B    ex_und          ; loop forever
ex_swi B    ex_swi          ; loop forever
ex_dabt B    ex_dabt        ; loop forever
ex_pabt B    ex_pabt        ; loop forever
int_irq B    int_irq        ; loop forever
int_fiq B    int_fiq        ; loop forever
```


Step 1: Take the Reset Exception

- Dummy handlers are set up
- Control is passed to code to initialize the hardware

Step 2: Start Initializing the Hardware

```
sandstone_init1
    LDR    r3, =SYSCFG      ; where SYSCFG=0x03ff0000
    LDR    r4, =0x03ffffa0
    STR    r4, [r3]
```

Address range	Size	Description
0x00000000 to 0x0003FFFF	256KB	32 bit SRAM bank, using ROMCON1
0x00040000 to 0x0007FFFF	256KB	32 bit SRAM bank, using ROMCON2
0x01800000 to 0x0187FFFF	512KB	16 bit flash bank, using ROMCON0
0x03FE0000 to 0x03FE1FFF	8KB	32 bit internal SRAM
0x03FF0000 to 0x03FFFFFF	64KB	Microcontroller register space



Step 2: Start Initializing the Hardware

```
; -- Section 2 : Initialize the hardware to use the segment display.
```

```
    LDR    r2, =SEG_MASK
    LDR    r3, =IOPMOD
    LDR    r4, [r3]
    ORR    r4, r4, r2
    STR    r4, [r3]
```

```
    LDR    r3, =IOPDATA
    LDR    r4, [r3]
    ORR    r4, r4, r2
    STR    r4, [r3]
```

```
    LDR    r2, =nSEG_MASK
```

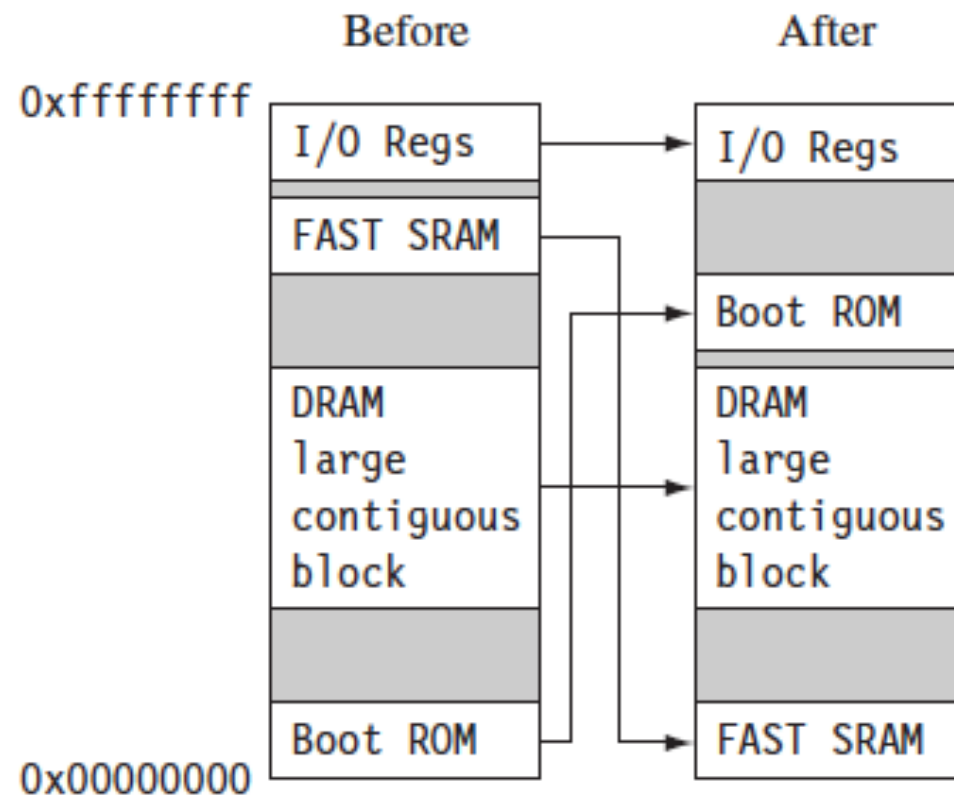
```
; -- To show hardware is configured correctly display 0 on
; -- the segment display.
```

```
    LDR    r5, data_dis0
    BL     _e7t_setSegmentDisplay
```

Step 2: Start Initializing the Hardware

- The system registers are set from a known base address—0x03ff0000
- The segment display is configured, so that it can be used to display progress

Step 3: Remap Memory



Step 3: Remap Memory

Memory type	Start address	End address	Size
Flash ROM	0x00000000	0x00080000	512K
SRAM bank 0	Unavailable	unavailable	256K
SRAM bank 1	Unavailable	unavailable	256K

```
LDR    lr, =sandstone_init2
LDR    r4, =0x1800000
ADD    lr, lr, r4

; -- Load in the target values into the
control registers

ADRL   r0, memorymaptable_str
LDMIA  r0, {r1-r12}
LDR    r0, =EXTDBWTH

; -- remap and jump to ROM code

STMIA  r0, {r1-r12}
MOV    pc, lr
```

Step 3: Remap Memory

```
memorymaptable_str
    DCD rEXTDBWTH      ; ROM0 (Half), ROM1 (Word), ROM1 (Word), rest Disabled
    DCD rROMCON0       ; 0x1800000 ~ 0x1880000, RCS0, 4Mbit, 4cycle
    DCD rROMCON1       ; 0x0000000 ~ 0x0040000, RCS1, 256KB, 2cycle
    DCD rROMCON2       ; 0x0040000 ~ 0x0080000, RCS2, 256KB, 2cycle
```

```
; -- ROMCON0 : ROM Bank0 Control register .....

ROMBasePtr0    EQU    0x180:SHL:10    ;=0x1800000
ROMEndPtr0     EQU    0x188:SHL:20    ;=0x1880000
PMC0           EQU    0x0             ; 0x0=Normal ROM, 0x1=4Word Page
                                           ; 0x2=8Word Page, 0x3=16Word Page
rTpa0          EQU    (0x0:SHL:2)     ; 0x0=5Cycle, 0x1=2Cycle
                                           ; 0x2=3Cycle, 0x3=4Cycle
rTacc0         EQU    (0x3:SHL:4)     ; 0x0=Disable, 0x1=2Cycle
                                           ; 0x2=3Cycle, 0x3=4Cycle
                                           ; 0x4=5Cycle, 0x5=6Cycle
                                           ; 0x6=7Cycle, 0x7=Reserved
rROMCON0       EQU    ROMEndPtr0+ROMBasePtr0+rTacc0+rTpa0+PMC0
```


Step 3: Remap Memory

Type	Start address	End address	Size
Flash ROM	0x01800000	0x01880000	512K
SRAM bank 0	0x00000000	0x00040000	256K
SRAM bank 1	0x00040000	0x00080000	256K

- Memory has been remapped
- pc now points to the next step. This address is located in the newly remapped flash ROM

Step 4: Initialize Communication Hardware

```
sandstone_init2  
  
        LDR        r5,data_dis4  
        BL         _e7t_setSegmentDisplay  
  
        B          sandstone_serialcode
```

- Serial port initialized—9600 baud, no parity, one stop bit, and no flow control
- Sandstone banner sent out through the serial port

Step 5: Bootloader – Copy Payload and Relinquish Control

```
sandstone_load_and_boot
; -- Section 1 : Copy payload to address 0x00000000
    MOV        r13,#0
    LDR        r12,payload_start_address
    LDR        r14,payload_end_address
block_copy
    LDMIA      r12!,{r0-r11}
    STMIA      r13!,{r0-r11}
    CMP        r12,r14
    BLE        block_copy
; -- Section 2 : Relinquish control over to the payload
    MOV        pc,#0
    LTORG
; -- Data Section : Payload information inserted at build time.
payload_start_address
    DCD        startAddress
payload_end_address
    DCD        endAddress
```

Step 5: Bootloader – Copy Payload and Relinquish Control

- Payload copied into SRAM, address 0x00000000
- Control of the pc is relinquished to the payload; pc = 0x00000000
- The system is completely booted


Minimal Operating System

Fundamental Components

- Initialization
- Memory handling (static memory management)
- Interrupt/exception handling
- Scheduler and context switch
- Device driver framework

Memory Map

- Non-MMU and Non-MPU



Step 1: Initiation

```
LDR      pc,vectorReset
        LDR      pc,vectorUndefined
        LDR      pc,vectorSWI
        LDR      pc,vectorPrefetchAbort
        LDR      pc,vectorDataAbort
        LDR      pc,vectorReserved
        LDR      pc,vectorIRQ
        LDR      pc,vectorFIQ

; -- Kernel Jump table -----
vectorReset      DCD      coreInitialize
vectorUndefined  DCD      coreUndefinedHandler
vectorSWI        DCD      coreSWIHandler
vectorPrefetchAbort DCD      corePrefetchAbortHandler
vectorDataAbort  DCD      coreDataAbortHandler
vectorReserved   DCD      coreReservedHandler
vectorIRQ        DCD      coreIRQHandler
vectorFIQ        DCD      coreFIQHandler

        END
```


Step 1: Initiation (Cont.)

```
coreInitialize
; -----
; Setup stacks for SVC,IRQ,SYS&USER
; -----
    MOV     sp,#0x80000          ; initializing sp to TopOfMemory
    MSR     CPSR_c,#NoInt|SYS32md
    MOV     sp,#0x40000          ; SYS/User Stack = 0x40000
    MSR     CPSR_c,#NoInt|IRQ32md
    MOV     sp,#0x9000          ; IRQ Stack = 0x9000
    MSR     CPSR_c,#NoInt|SVC32md
; -----
; Setup Task Process Control Block (PCB).
; -----
    LDR     r0,=C_EntryTask2      ; addr of C_EntryTask2
    LDR     r1,=PCB_PtrTask2
    MOV     r2,#0x10000
    BL      pcbSetUp
; -- set the current ID to TASK1 .....
    LDR     r0, =PCB_CurrentTask
    MOV     r1, #0
    STR     r1,[r0]                ; first task ID = 0
    LDR     lr,=C_Entry
    MOV     pc,lr                  ; call C_Entry
END
```

Step 1: Initiation (Cont.)

```
int C_Entry(void)
{
    lltrace(cinit_init(), CINITINIT);
    lltrace(eventTickStart(), TICKSTART);
    STATE=BOOT_SLOS;
    __asm
    {
        MSR                CPSR_c, #0x50
    }
    lltrace(C_EntryTask1(), ENTERTASK);
    return 0;
}
```

Step 1: Initiation (Cont.)

```
void C_EntryTask1(void)
{
volatile int delay;
    if (openRedLED ()==0)
    {
        while (1)
        {
            aWAIT;
            switchOnRedLED ();
            /* dummy time delay... */
            for (delay=0; delay<0x20ffff; delay++) {}
            aSIGNAL;
            switchOffRedLED ();
            /* dummy time delay... */
            for (delay=0; delay<0x20ffff; delay++) {}
        }
    }
    while (1)
    {
        delay=0xBEEFBEEF;
    };
}
```

Step 3: Interrupt and Exception Handling

Exception	Purpose
Reset	initialize the operating system
SWI	mechanism to access device drivers
IRQ	mechanism to service events

Step 3: Interrupt and Exception Handling (Cont.)

coreSWIHandler

```
STMFD    sp!,{r0-r12,r14}    ; save context
LDR      r10,[r14,#-4]        ; load SWI instruction
BIC      r10,r10,#0xff000000  ; mask off the MSB 8 bits
MOV      r1,r13               ; copy r13_svc to r1
MRS      r2,spsr              ; copy spsr to r2
STMFD    r13!,{r2}           ; save r2 onto the stack
BL       swi_jumtable         ; branch to the swi_jumtable
```

```
LDMFD    r13!,{r2}           ; restore the r2 (spsr)
MSR      spsr_cxsf,r2        ; copy r2 back to spsr
LDMFD    r13!,{r0-r12,pc}^    ; restore context and return
```

swi_jumtable

```
MOV      r0,r10               ; move the SWI number to r0
B        eventsSWIHandler     ; branch to SWI handler
```

```
void eventsSWIHandler(int swi_number, SwiRegs *r)
{
    if (swi_number==SLOS)
    {
        if (r->r[0]==Event_IODeviceInit)
        {
            /* do not enable IRQ interrupts .... */
            io_initialize_drivers ();
        }
        else
        {
            /* if not initializing change to system mode
               and enable IRQs */
            if (STATE!=1) {modifyControlCPSR (SYSTEM|IRQoN);}

            switch (r->r[0])
            {
            case /* SWI */ Event_IODeviceOpen:
                r->r[0] =
                    (unsigned int) io_open_driver
                    (
                        /*int *ID */ (UID *)r->r[1],
                        /*unsigned major_device */ r->r[2],
                        /*unsigned minor_device */ r->r[3]
                    );
            }
        }
    }
}
```

Step 3: Interrupt and Exception Handling (Cont.)

```
TICKINT    EQU 0x400
BUTTONINT  EQU 0x001
```

```
eventsIRQHandler
```

```
    SUB      r14, r14, #4           ; r14_irq-=4
    STMFD    r13!, {r0-r3, r12, r14} ; save context
    LDR      r0,INTPND             ; r0=int pending reg
    LDR      r0,[r0]               ; r0=memory[r0]
    TST      r0,#TICKINT           ; if tick int
    BNE      eventsTickVeneer      ; then tick ISR
    TST      r0,#BUTTONINT         ; if button interrupt
    BNE      eventsButtonVeneer    ; then button ISR
    LDMFD    r13!, {r0-r3, r12, pc}^ ; return to task
```

```
eventsTickVeneer
```

```
    BL      eventsTickService      ; reset tick hardware
    B       kernelScheduler        ; branch to scheduler
```

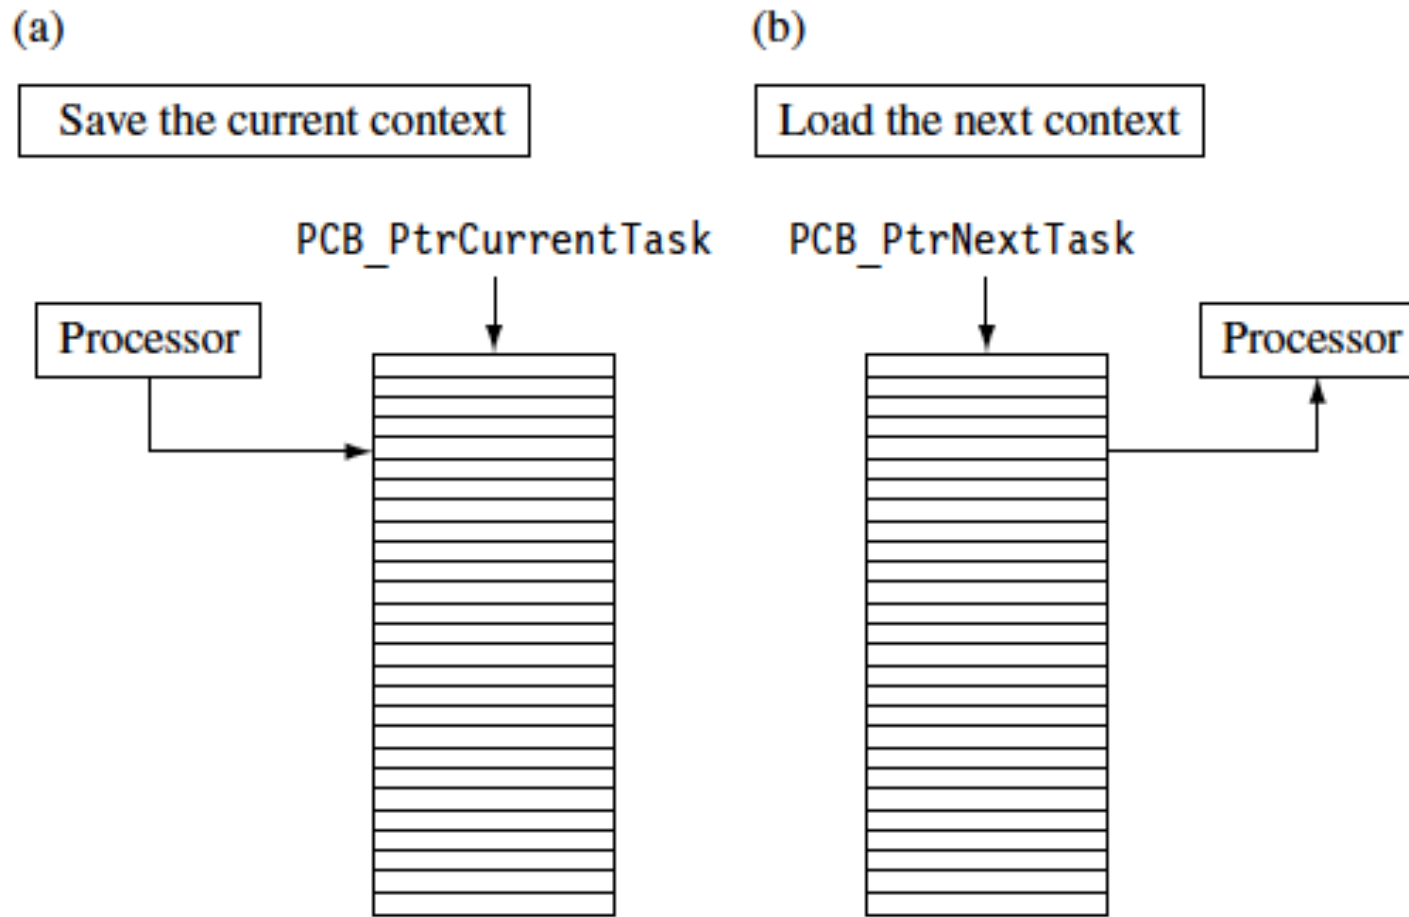
Step 4: Scheduler and Context Switch

```
task t=0,t';

scheduler()
{
    t' = t + 1;
    if t' = MAX_NUMBER_OF_TASKS then
        t' = 0 // the first task.
    end;

    ContextSwitch(t,t')
}
```


Step 4: Scheduler and Context Switch (Cont.)



Step 4: Scheduler and Context Switch (Cont.)

- Saving registers to a PCB

```
Offset15Regs EQU 15*4
```

```
handler_contextswitch
```

```
    LDMFD    r13!,{r0-r3,r12,r14}    ; [1.1] restore registers
    LDR      r13,=PCB_PtrCurrentTask ; [1.2]
    LDR      r13,[r13]                ; r13=mem32[r13]
    SUB      r13,r13,#Offset15Regs    ; r13-=15*Reg:place r13
    STMIA    r13,{r0-r14}^            ; [1.3] save user mode registers
    MRS      r0, spsr                 ; copy spsr
    STMDB    r13,{r0,r14}             ; save r0(spsr) & r14(lr)
```

Step 4: Scheduler and Context Switch (Cont.)

- Loading registers from a PCB

```
LDR    r13,=PCB_PtrNextTask ; [2.1] r13=PCB_PtrNextTask
LDR    r13,[r13]             ; r13=mem32[r13] : next PCB
SUB    r13,r13,#Offset15Regs ; r13-=15*Registers
LDMDB  r13,{r0,r14}          ; [2.2] load r0 & r14
MSR    spsr_cxsf, r0         ; spsr = r0
LDMIA  r13,{r0-r14}^         ; load r0_user-r14_user
LDR    r13,=PCB_IRQStack     ; [2.3] r13=IRQ stack addr
LDR    r13,[r13]             ; r13=mem32[r13] : reset IRQ
MOVS   pc,r14                ; [2.4] return to next task
```

Step 5: Device Driver Framework

```
host = eventIODeviceOpen(&serial,DEVICE_SERIAL_E7T,COM1);

if (host==0)
{
    /* ...error device driver not found...*/
}

switch (serial)
{
case DEVICE_IN_USE:
case DEVICE_UNKNOWN:
    /* ...problem with device... */
}
```

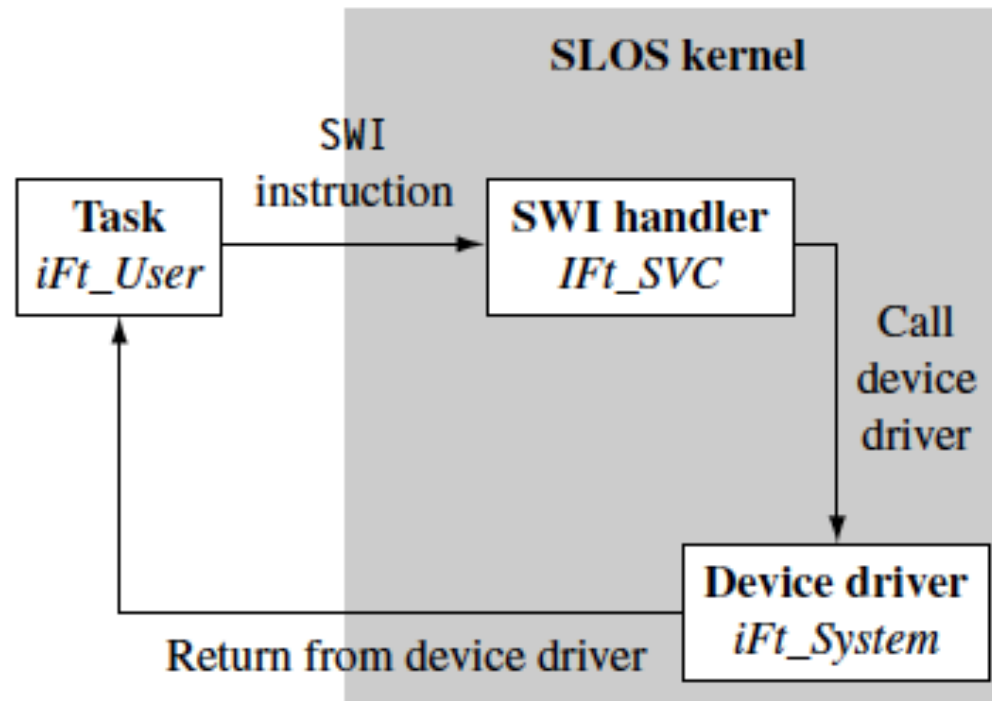
Step 5: Device Driver Framework (Cont.)

```
PRE    r0 = Event_IODeviceOpen (unsigned int)
        r1 = &serial (UID *u)
        r2 = DEVICE_SERIAL_E7T (unsigned int major)
        r3 = COM1 (unsigned int minor)

        SWI      5075

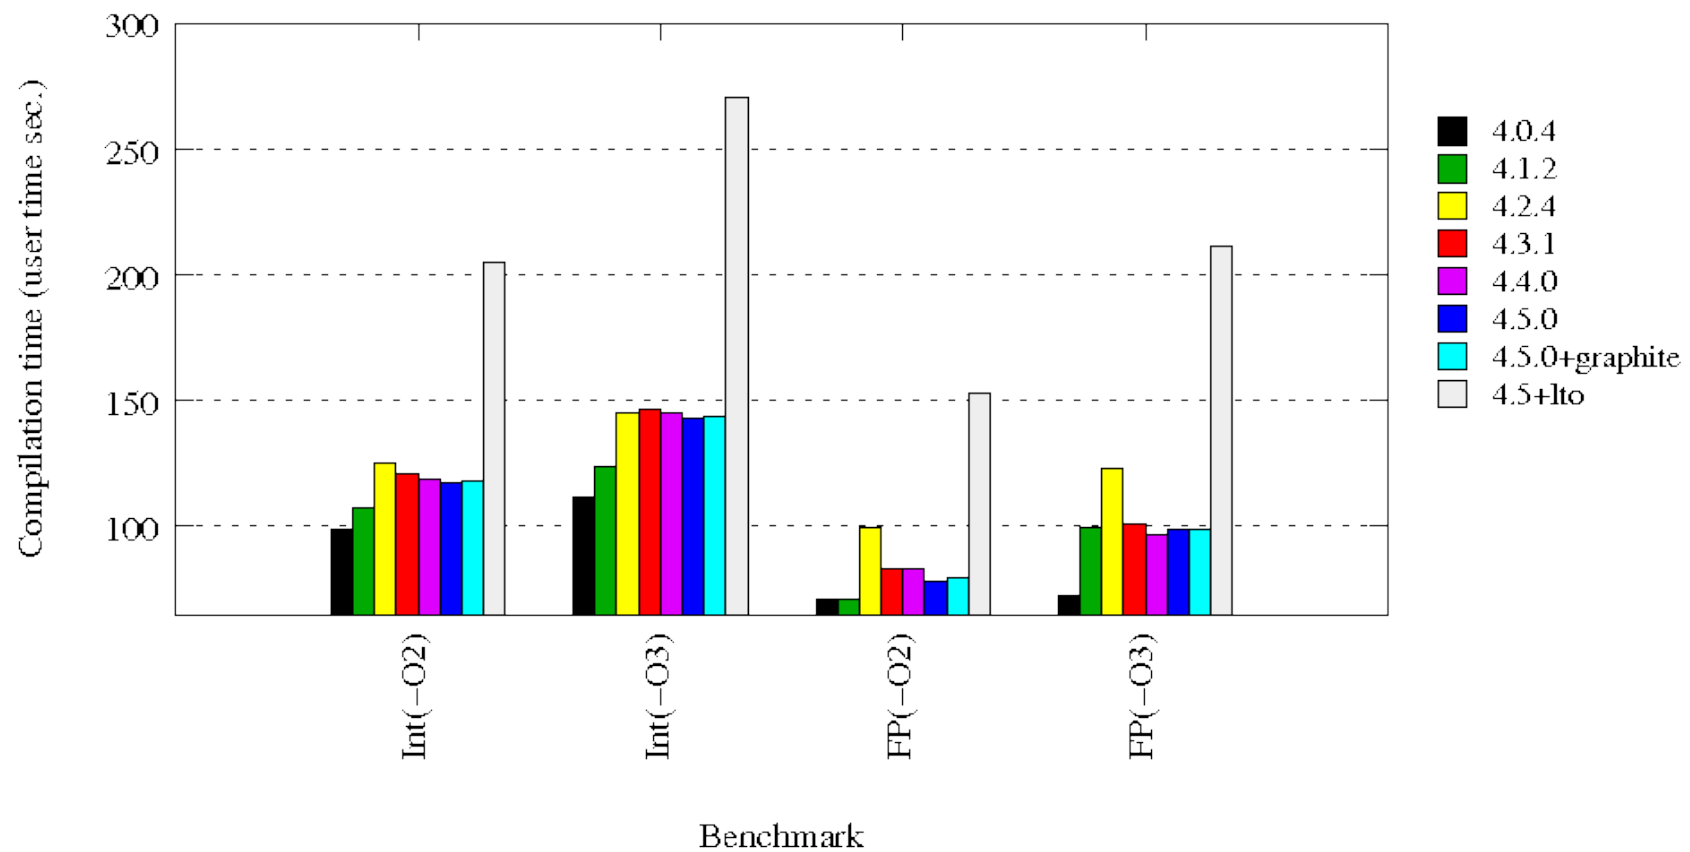
POST    r1 = The data pointed to by the UID pointer is updated
```

Step 5: Device Driver Framework (Cont.)

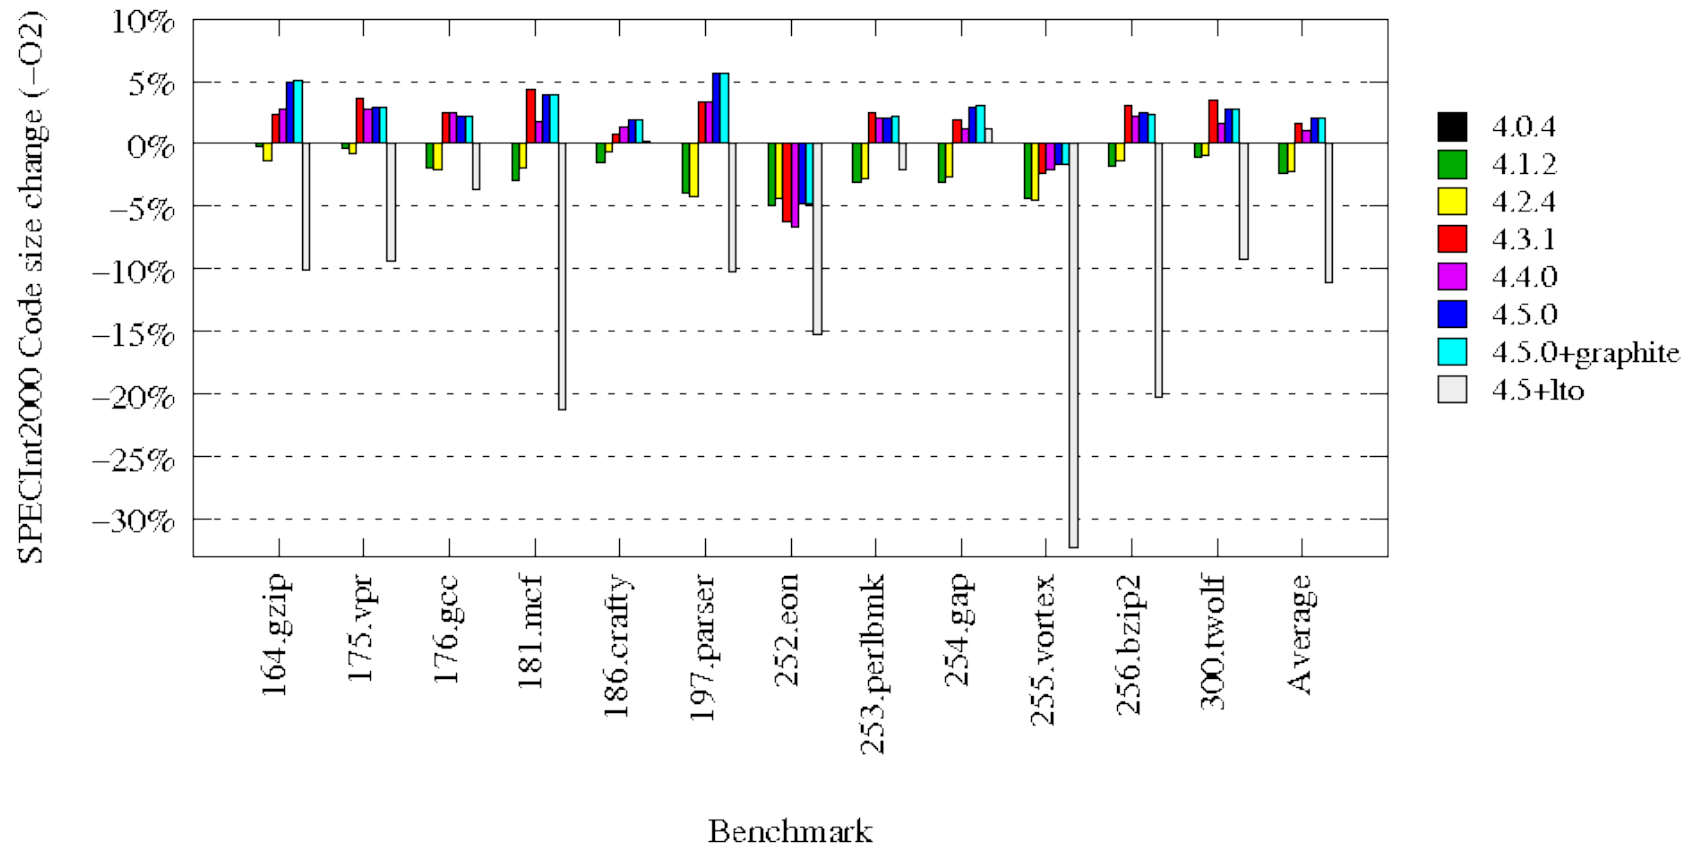


Efficient C Programming


Compiler Versions



Compiler Versions (Cont.)



Compiler Versions (Cont.)



Compiler Options

- -O0
 - Reduce compilation time and make debugging produce the expected results. This is the default
- -O1
 - Optimizing compilation takes somewhat more time, and a lot more memory for a large function
- -O2
 - Optimize even more. GCC performs nearly all supported optimizations that do not involve a space-speed tradeoff
- -O3
 - turns on all optimizations specified by -O2 and also turns on the -finline-functions, -funswitch-loops, -fpredictive-commoning, ...

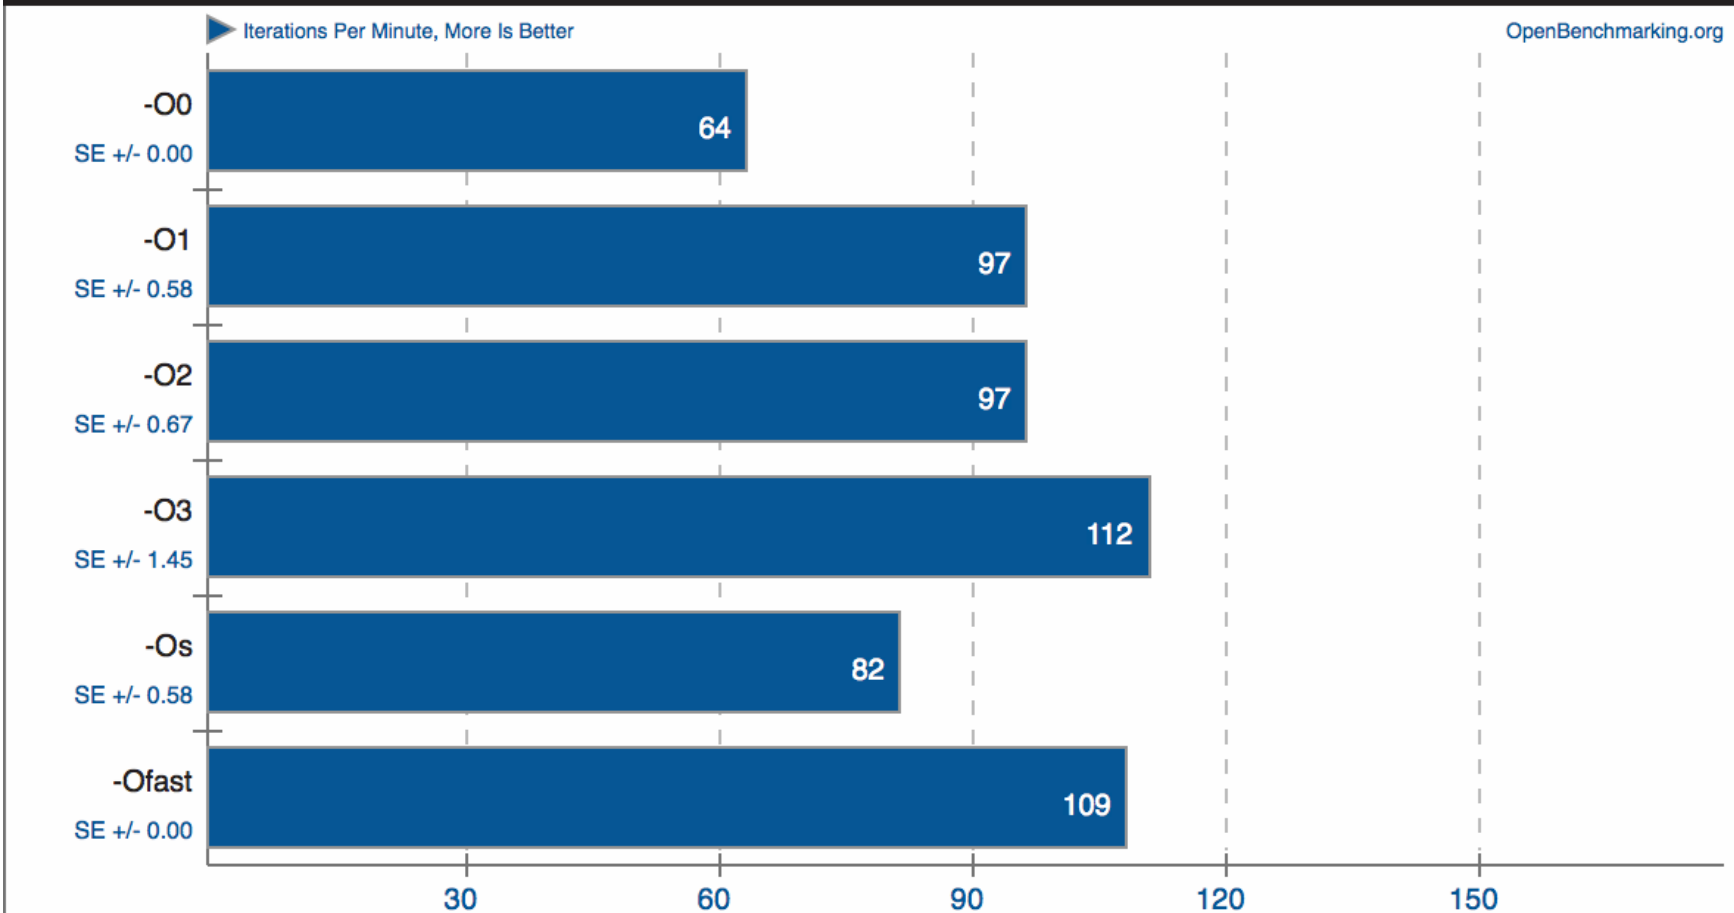
Compiler Options (Cont.)

- `Os`
 - Optimize for size. `-Os` enables all `-O2` optimizations that do not typically increase code size
- `-Ofast`
 - enables all `-O3` optimizations. It also enables optimizations that are not valid for all standard-compliant programs

Performance

GraphicsMagick v1.3.12

Operation: Blur




Powered By Phoronix Test Suite 4.2.0m0

1. (CC) gcc options: -std=gnu99 -fopenmp -pthread -lz -lm -lgomp -lpthread

Compilation Time

Timed PHP Compilation v5.2.9

Time To Compile



1. (CC) gcc options: -pedantic -ldl -lz -lm

Powered By Phoronix Test Suite 4.2.0m0

Examples

- What are the problems?

```
void memclr(char *data, int N)
{
    for (; N>0; N--)
    {
        *data=0;
        data++;
    }
}
```

Technologies for Better C Code

- Basic C Data Types
- C Looping Structures
- Register Allocation
- Function Calls
- Pointer Aliasing
- Structure Arrangement
- Bit-fields
- Unaligned Data and Endianness
- Division
- Floating Point
- Inline Functions and Inline Assembly

Local Variable Types

<code>int checksum_v1(int *data)</code>	<code>checksum_v1</code>	
<code>{</code>	<code>MOV</code>	<code>r2,r0 ; r2 = data</code>
<code>char i;</code>	<code>MOV</code>	<code>r0,#0 ; sum = 0</code>
<code>int sum=0;</code>	<code>MOV</code>	<code>r1,#0 ; i = 0</code>
	<code>checksum_v1_loop</code>	
<code>for (i=0; i<64; i++)</code>	<code>LDR</code>	<code>r3,[r2,r1,LSL #2] ; r3 = data[i]</code>
<code>{</code>	<code>ADD</code>	<code>r1,r1,#1 ; r1 = i+1</code>
<code>sum += data[i];</code>	<code>AND</code>	<code>r1,r1,#0xff ; i = (char)r1</code>
<code>}</code>	<code>CMP</code>	<code>r1,#0x40 ; compare i, 64</code>
<code>return sum;</code>	<code>ADD</code>	<code>r0,r3,r0 ; sum += r3</code>
<code>}</code>	<code>BCC</code>	<code>checksum_v1_loop ; if (i<64) loop</code>
	<code>MOV</code>	<code>pc,r14 ; return sum</code>

Local Variable Types

```
checksum_v2
    MOV    r2,r0          ; r2 = data
    MOV    r0,#0          ; sum = 0
    MOV    r1,#0          ; i = 0
checksum_v2_loop
    LDR    r3,[r2,r1,LSL #2] ; r3 = data[i]
    ADD    r1,r1,#1        ; r1++
    CMP    r1,#0x40        ; compare i, 64
    ADD    r0,r3,r0        ; sum += r3
    BCC    checksum_v2_loop ; if (i<64) goto loop
    MOV    pc,r14          ; return sum
```

Local Variable Types

	checksum_v3	
short checksum_v3(short *data)	MOV r2,r0	; r2 = data
{	MOV r0,#0	; sum = 0
unsigned int i;	MOV r1,#0	; i = 0
short sum=0;	checksum_v3_loop	
	ADD r3,r2,r1,LSL #1	; r3 = &data[i]
for (i=0; i<64; i++)	LDRH r3,[r3,#0]	; r3 = data[i]
{	ADD r1,r1,#1	; i++
sum = (short)(sum + data[i]);	CMP r1,#0x40	; compare i, 64
}	ADD r0,r3,r0	; r0 = sum + r3
return sum;	MOV r0,r0,LSL #16	
}	MOV r0,r0,ASR #16	; sum = (short)r0
	BCC checksum_v3_loop	; if (i<64) goto loop
	MOV pc,r14	; return sum

Local Variable Types

```
short checksum_v4(short *data)
{
    unsigned int i;
    int sum=0;

    for (i=0; i<64; i++)
    {
        sum += *(data++);
    }
    return (short)sum;
}
```

```
checksum_v4
    MOV     r2,#0           ; sum = 0
    MOV     r1,#0           ; i = 0
checksum_v4_loop
    LDRSH   r3,[r0],#2      ; r3 = *(data++)
    ADD     r1,r1,#1        ; i++
    CMP     r1,#0x40        ; compare i, 64
    ADD     r2,r3,r2        ; sum += r3
    BCC     checksum_v4_loop ; if (sum<64) goto loop
    MOV     r0,r2,LSL #16
    MOV     r0,r0,ASR #16   ; r0 = (short)sum
    MOV     pc,r14         ; return r0
```

Functional Argument Types

```
add_v1
    ADD    r0,r0,r1,ASR #1      ; r0 = (int)a + ((int)b>>1)
    MOV    r0,r0,LSL #16
    MOV    r0,r0,ASR #16       ; r0 = (short)r0
    MOV    pc,r14              ; return r0
```

```
short add_v1(short a, short b)
{
    return a + (b>>1);
}
```

```
add_v1_gcc
    MOV    r0, r0, LSL #16
    MOV    r1, r1, LSL #16
    MOV    r1, r1, ASR #17      ; r1 = (int)b>>1
    ADD    r1, r1, r0, ASR #16  ; r1 += (int)a
    MOV    r1, r1, LSL #16
    MOV    r0, r1, ASR #16     ; r0 = (short)r1
    MOV    pc, lr              ; return r0
```

The Efficient Use of C Types

- For local variables held in registers, don't use a char or short type unless 8-bit or 16-bit modular arithmetic is necessary. Use the signed or unsigned int types instead. Unsigned types are faster when you use divisions
- For array entries and global variables held in main memory, use the type with the smallest size possible to hold the required data. This saves memory footprint

The Efficient Use of C Types (Cont.)

- Use explicit casts when reading array entries or global variables into local variables, or writing local variables out to array entries
- Avoid implicit or explicit narrowing casts in expressions because they usually cost extra cycles
- Avoid char and short types for function arguments or return values. Instead use the int type even if the range of the parameter is smaller. This prevents the compiler performing unnecessary casts

Loops with a Fixed Number of Iterations

```
int checksum_v5(int *data)
{
    unsigned int i;
    int sum=0;

    for (i=0; i<64; i++)
    {
        sum += *(data++);
    }
    return sum;
}
```

```
checksum_v5
    MOV     r2,r0           ; r2 = data
    MOV     r0,#0           ; sum = 0
    MOV     r1,#0           ; i = 0
checksum_v5_loop
    LDR     r3,[r2],#4       ; r3 = *(data++)
    ADD     r1,r1,#1         ; i++
    CMP     r1,#0x40         ; compare i, 64
    ADD     r0,r3,r0         ; sum += r3
    BCC     checksum_v5_loop ; if (i<64) goto loop
    MOV     pc,r14           ; return sum
```


Loops with a Fixed Number of Iterations

```
int checksum_v6(int *data)
{
    unsigned int i;
    int sum=0;

    for (i=64; i!=0; i--)
    {
        sum += *(data++);
    }
    return sum;
}
```

```
checksum_v6
    MOV     r2,r0           ; r2 = data
    MOV     r0,#0           ; sum = 0
    MOV     r1,#0x40        ; i = 64
checksum_v6_loop
    LDR     r3,[r2],#4       ; r3 = *(data++)
    SUBS    r1,r1,#1         ; i-- and set flags
    ADD     r0,r3,r0         ; sum += r3
    BNE     checksum_v6_loop ; if (i!=0) goto loop
    MOV     pc,r14           ; return sum
```

Loops Using A Variable Number of Iterations

```
int checksum_v7(int *data, unsigned int N)
{
    int sum=0;

    for (; N!=0; N--)
    {
        sum += *(data++);
    }
    return sum;
}
```

checksum_v7

```
MOV    r2,#0                ; sum = 0
CMP    r1,#0                ; compare N, 0
BEQ    checksum_v7_end     ; if (N==0) goto end

checksum_v7_loop
LDR    r3,[r0],#4           ; r3 = *(data++)
SUBS   r1,r1,#1             ; N-- and set flags
ADD    r2,r3,r2             ; sum += r3
BNE    checksum_v7_loop    ; if (N!=0) goto loop

checksum_v7_end
MOV    r0,r2                ; r0 = sum
MOV    pc,r14               ; return r0
```

Loops Using A Variable Number of Iterations

```
int checksum_v8(int *data, unsigned int N)
```

```
{
```

```
    int sum=0;
```

```
    do
```

```
    {
```

```
        sum += *(data++);
```

```
    } while (--N!=0);
```

```
    return sum;
```

```
}
```

```
checksum_v8
```

```
    MOV    r2,#0
```

```
; sum = 0
```

```
checksum_v8_loop
```

```
    LDR    r3,[r0],#4
```

```
; r3 = *(data++)
```

```
    SUBS   r1,r1,#1
```

```
; N-- and set flags
```

```
    ADD    r2,r3,r2
```

```
; sum += r3
```

```
    BNE    checksum_v8_loop
```

```
; if (N!=0) goto loop
```

```
    MOV    r0,r2
```

```
; r0 = sum
```

```
    MOV    pc,r14
```

```
; return r0
```

Loop Unrolling

```
int checksum_v9(int *data, unsigned int N)
{
    int sum=0;
do
{
    sum += *(data++);
    sum += *(data++);
    sum += *(data++);
    sum += *(data++);
    N -= 4;
} while ( N!=0);
return sum;
}
```

checksum_v9

```
MOV r2,#0      ; sum = 0
checksum_v9_loop
LDR r3,[r0],#4  ; r3 = *(data++)
SUBS r1,r1,#4   ; N -= 4 & set flags
ADD r2,r3,r2    ; sum += r3
LDR r3,[r0],#4  ; r3 = *(data++)
ADD r2,r3,r2    ; sum += r3
LDR r3,[r0],#4  ; r3 = *(data++)
ADD r2,r3,r2    ; sum += r3
LDR r3,[r0],#4  ; r3 = *(data++)
ADD r2,r3,r2    ; sum += r3
LDR r3,[r0],#4  ; r3 = *(data++)
ADD r2,r3,r2    ; sum += r3
BNE checksum_v9_loop ; if (N!=0) goto loop
MOV r0,r2       ; r0 = sum
MOV pc,r14      ; return r0
```

Writing Loops Efficiently

- Use loops that count down to zero. Then the compiler does not need to allocate a register to hold the termination value, and the comparison with zero is free
- Use unsigned loop counters by default and the continuation condition $i \neq 0$ rather than $i > 0$. This will ensure that the loop overhead is only two instructions
- Use do-while loops rather than for loops when you know the loop will iterate at least once. This saves the compiler checking to see if the loop count is zero

Writing Loops Efficiently (Cont.)

- Unroll important loops to reduce the loop overhead. Do not overunroll. If the loop overhead is small as a proportion of the total, then unrolling will increase code size and hurt the performance of the cache
- Try to arrange that the number of elements in arrays are multiples of four or eight. You can then unroll loops easily by two, four, or eight times without worrying about the leftover array elements

Register Allocations

- Try to limit the number of local variables in the internal loop of functions to 12. The compiler should be able to allocate these to ARM registers
- You can guide the compiler as to which variables are important by ensuring these variables are used within the innermost loop

Function Calls

...	...
$sp + 16$	Argument 8
$sp + 12$	Argument 7
$sp + 8$	Argument 6
$sp + 4$	Argument 5
sp	Argument 4

$r3$	Argument 3	
$r2$	Argument 2	
$r1$	Argument 1	
$r0$	Argument 0	Return value


```

char *queue_bytes_v1(
    char *Q_start,          /* Queue buffer start address */
    char *Q_end,            /* Queue buffer end address */
    char *Q_ptr,            /* Current queue pointer position */
    char *data,             /* Data to insert into the queue */
    unsigned int N)         /* Number of bytes to insert */
{
    do
    {
        *(Q_ptr++) = *(data++);

        if (Q_ptr == Q_end)
        {
            Q_ptr = Q_start;
        }
    } while (--N);
    return Q_ptr;
}

queue_bytes_v1
    STR    r14,[r13,#-4]!    ; save lr on the stack
    LDR    r12,[r13,#4]     ; r12 = N

queue_v1_loop
    LDRB   r14,[r3],#1      ; r14 = *(data++)
    STRB   r14,[r2],#1      ; *(Q_ptr++) = r14
    CMP    r2,r1            ; if (Q_ptr == Q_end)
    MOVEQ  r2,r0            ;     {Q_ptr = Q_start;}
    SUBS   r12,r12,#1       ; --N and set flags
    BNE    queue_v1_loop    ; if (N!=0) goto loop
    MOV    r0,r2            ; r0 = Q_ptr
    LDR    pc,[r13],#4      ; return r0

```


Calling Functions Efficiently

- Try to restrict functions to four arguments. This will make them more efficient to call. Use structures to group related arguments and pass structure pointers instead of multiple arguments
- Define small functions in the same source file and before the functions that call them. The compiler can then optimize the function call or inline the small function
- Critical functions can be inlined using the `__inline` keyword

Pointer Aliasing

```
void timers_v1(int *timer1, int *timer2, int *step)
{
    *timer1 += *step;
    *timer2 += *step;
}
```

```
timers_v1
    LDR    r3,[r0,#0]        ; r3 = *timer1
    LDR    r12,[r2,#0]       ; r12 = *step
    ADD    r3,r3,r12         ; r3 += r12
    STR    r3,[r0,#0]        ; *timer1 = r3
    LDR    r0,[r1,#0]        ; r0 = *timer2
    LDR    r2,[r2,#0]        ; r2 = *step
    ADD    r0,r0,r2          ; r0 += r2
    STR    r0,[r1,#0]        ; *timer2 = t0
    MOV    pc,r14            ; return
```

Pointer Aliasing (Cont.)

```
typedef struct {int step;} State;  
typedef struct {int timer1, timer2;} Timers;  
  
void timers_v2(State *state, Timers *timers)  
{  
    timers->timer1 += state->step;  
    timers->timer2 += state->step;  
}
```

```
void timers_v3(State *state, Timers *timers)  
{  
    int step = state->step;  
  
    timers->timer1 += step;  
    timers->timer2 += step;  
}
```

Avoiding Pointer Aliasing

- Do not rely on the compiler to eliminate common subexpressions involving memory accesses. Instead create new local variables to hold the expression. This ensures the expression is evaluated only once
- Avoid taking the address of local variables. The variable may be inefficient to access from then on

Structure Arrangement

```
struct {  
    char a;  
    int b;  
    char c;  
    short d;  
}
```

Address	+3	+2	+1	+0
+0	pad	pad	pad	a
+4	b[31,24]	b[23,16]	b[15,8]	b[7,0]
+8	d[15,8]	d[7,0]	pad	c

Structure Arrangement (Cont.)

```
struct {  
    char a;  
    char c;  
    short d;  
    int b;  
}
```

Address	+3	+2	+1	+0
+0	d[15,8]	d[7,0]	c	a
+4	b[31,24]	b[23,16]	b[15,8]	b[7,0]

Structure Arrangement (Cont.)

```
__packed struct {  
    char a;  
    int b;  
    char c;  
    short d;  
}
```

Address	+3	+2	+1	+0
+0	b[23,16]	b[15,8]	b[7,0]	a
+4	d[15,8]	d[7,0]	c	b[31,24]

Efficient Structure Arrangement

- Lay structures out in order of increasing element size. Start the structure with the smallest elements and finish with the largest
- Avoid very large structures. Instead use a hierarchy of smaller structures
- For portability, manually add padding (that would appear implicitly) into API structures so that the layout of the structure does not depend on the compiler
- Beware of using enum types in API structures. The size of an enum type is compiler dependent

Bit Fields

```
void dostageA(void);  
void dostageB(void);  
void dostageC(void);
```

```
typedef struct {  
    unsigned int stageA : 1;  
    unsigned int stageB : 1;  
    unsigned int stageC : 1;  
} Stages_v1;
```

```
void dostages_v1(Stages_v1 *stages)  
{  
    if (stages->stageA)  
    {  
        dostageA();  
    }  
}
```

```
    if (stages->stageB)  
    {  
        dostageB();  
    }  
    if (stages->stageC)  
    {  
        dostageC();  
    }  
}
```

Bit Fields (Cont.)

dostages_v1

```
STMFD    r13!,{r4,r14}    ; stack r4, lr
MOV      r4,r0             ; move stages to r4
LDR      r0,[r0,#0]        ; r0 = stages bitfield
TST      r0,#1             ; if (stages->stageA)
BLNE     dostageA          ;     {dostageA();}
LDR      r0,[r4,#0]        ; r0 = stages bitfield
MOV      r0,r0,LSL #30     ; shift bit 1 to bit 31
CMP      r0,#0             ; if (bit31)
BLLT     dostageB          ;     {dostageB();}
LDR      r0,[r4,#0]        ; r0 = stages bitfield
MOV      r0,r0,LSL #29     ; shift bit 2 to bit 31
CMP      r0,#0             ; if (!bit31)
LDMLTFD  r13!,{r4,r14}    ;     return
BLT      dostageC          ; dostageC();
LDMFD    r13!,{r4,pc}     ; return
```

Bit Fields (Cont.)

```
typedef unsigned long Stages_v2;  
#define STAGEA (1ul << 0)
```

```
#define STAGEB (1ul << 1)  
#define STAGEC (1ul << 2)  
  
void dostages_v2(Stages_v2 *stages_v2)  
{  
    Stages_v2 stages = *stages_v2;  
  
    if (stages & STAGEA)  
    {  
        dostageA();  
    }  
    if (stages & STAGEB)  
    {  
        dostageB();  
    }  
    if (stages & STAGEC)  
    {  
        dostageC();  
    }  
}
```

Bit Fields (Cont.)

```
dostages_v2
    STMFD    r13!,{r4,r14}    ; stack r4, lr
    LDR      r4,[r0,#0]       ; stages = *stages_v2
    TST      r4,#1            ; if (stage & STAGEA)
    BLNE     dostageA         ;     {dostageA();}
    TST      r4,#2            ; if (stage & STAGEB)
    BLNE     dostageB         ;     {dostageB();}
    TST      r4,#4            ; if (!(stage & STAGEC))
    LDMNEFD  r13!,{r4,r14}    ; return;
    BNE      dostageC         ; dostageC();
    LDMFD    r13!,{r4,pc}     ; return
```

Bit-fields

- Avoid using bit-fields. Instead use `#define` or `enum` to define mask values
- Test, toggle, and set bit-fields using integer logical AND, OR, and exclusive OR operations with the mask values. These operations compile efficiently, and you can test, toggle, or set multiple fields at the same time

Unaligned Data and Endiannes

Little-endian configuration.

Instruction	Width (bits)	b31..b24	b23..b16	b15..b8	b7..b0
LDRB	8	0	0	0	B(A)
LDRSB	8	S(A)	S(A)	S(A)	B(A)
STRB	8	X	X	X	B(A)
LDRH	16	0	0	B(A+1)	B(A)
LDRSH	16	S(A+1)	S(A+1)	B(A+1)	B(A)
STRH	16	X	X	B(A+1)	B(A)
LDR/STR	32	B(A+3)	B(A+2)	B(A+1)	B(A)

Unaligned Data and Endiannes (Cont.)

Big-endian configuration.

Instruction	Width (bits)	b31..b24	b23..b16	b15..b8	b7..b0
LDRB	8	0	0	0	B(A)
LDRSB	8	S(A)	S(A)	S(A)	B(A)
STRB	8	X	X	X	B(A)
LDRH	16	0	0	B(A)	B(A+1)
LDRSH	16	S(A)	S(A)	B(A)	B(A+1)
STRH	16	X	X	B(A)	B(A+1)
LDR/STR	32	B(A)	B(A+1)	B(A+2)	B(A+3)

Endianness and Alignment

- Avoid using unaligned data if you can
- Use the type `char *` for data that can be at any byte alignment. Access the data by reading bytes and combining with logical operations. Then the code won't depend on alignment or ARM endianness configuration
- For fast access to unaligned structures, write different variants according to pointer alignment and processor endianness

Division

```
offset = (offset + increment) % buffer_size;
```

```
offset += increment;  
if (offset >= buffer_size)  
{  
    offset -= buffer_size;  
}
```

Converting Divides into Multiples

```
void scale(  
    unsigned int *dest,          /* destination for the scale data */  
    unsigned int *src,          /* source unscaled data */  
    unsigned int d,             /* denominator to divide by */  
    unsigned int N)             /* data length */  
{  
    unsigned int s = 0xFFFFFFFFu / d;
```

Converting Divides into Multiples (Cont.)

```
do
{
    unsigned int n, q, r;

    n = *(src++);
    q = (unsigned int)(((unsigned long long)n * s) >> 32);
    r = n - q * d;
    if (r >= d)
    {
        q++;
    }
    *(dest++) = q;
} while (--N);
}
```

Unsigned Division by a Constant

```
unsigned int udiv_by_const(unsigned int n, unsigned int d)
{
    unsigned int s,k,q;

    /* We assume d!=0 */

    /* first find k such that  $(1 \ll k) \leq d < (1 \ll (k+1))$  */
    for (k=0; d/2>=(1u<<k); k++);

    if (d==1u<<k)
    {
        /* we can implement the divide with a shift */
        return n>>k;
    }

    /* d is in the range  $(1 \ll k) < d < (1 \ll (k+1))$  */
    s = (unsigned int)(((1ull<<(32+k))+(1ull<<k))/d);
```

Unsigned Division by a Constant (Cont.)

```
if ((unsigned long long)s*d >= (1ull << (32+k)))
{
    /* n/d = (n*s) >> (32+k) */
    q = (unsigned int)((((unsigned long long)n*s) >> 32));
    return q >> k;
}

/* n/d = (n*s+s) >> (32+k) */

q = (unsigned int)((((unsigned long long)n*s + s) >> 32));
return q >> k;
}
```

Division

- Avoid divisions as much as possible. Do not use them for circular buffer handling
- If you can't avoid a division, then try to take advantage of the fact that divide routines often generate the quotient n/d and modulus $n\%d$ together
- To repeatedly divide by the same denominator d , calculate $s = (2^k - 1)/d$ in advance. You can replace the divide of a k -bit unsigned integer by d with a $2k$ -bit multiply by s

Division (Cont.)

- To divide unsigned $n < 2^N$ by an unsigned constant d , you can find a 32-bit unsigned s and shift k such that n/d is either $(ns) \gg (N + k)$ or $(ns + s) \gg (N + k)$. The choice depends only on d . There is a similar result for signed divisions

Floating Point

- The majority of ARM processor implementations do not provide hardware floating-point support, which saves on power and area
- In practice, this means that the C compiler converts every floating-point operation into a subroutine call. The C library contains subroutines to simulate floating-point behavior using integer arithmetic. This code is written in highly optimized assembly
- If you need fast execution and fractional values, you should use fixed-point or blockfloating algorithms

Inline Functions and Inline Assembly

- Use inline functions to declare new operations or primitives not supported by the C compiler
- Use inline assembly to access ARM instructions not supported by the C compiler. Examples are coprocessor instructions or ARMv5E extensions.

Efficient Assembly Programming

Technologies to Write Efficient Assembly Code

- Instruction Scheduling
- Register Allocation
- Conditional Execution
- Bit Manipulation
- Efficient Switches
- Handling Unaligned Data

Write Assembly Code

```
#include <stdio.h>

int square(int i);

int main(void)
{
    int i;

    for (i=0; i<10; i++)
    {
        printf("Square of %d is %d\n", i, square(i));
    }
}

int square(int i)
{
    return i*i;
}
```

Write Assembly Code (Cont.)

```
AREA    |.text|, CODE, READONLY

EXPORT  square

; int square(int i)
square
    MUL    r1, r0, r0    ; r1 = r0 * r0
    MOV    r0, r1        ; r0 = r1
    MOV    pc, lr        ; return r0
END
```

Write Assembly Code (Cont.)

```
AREA    |.text|, CODE, READONLY

EXPORT  main

IMPORT  |Lib$$Request$$armlib|, WEAK
IMPORT  __main      ; C library entry
IMPORT  printf      ; prints to stdout

i      RN 4

; int main(void)

main
    STMFD    sp!, {i, lr}
    MOV      i, #0
```


Write Assembly Code (Cont.)

```
loop
    ADR    r0, print_string
    MOV    r1, i
    MUL    r2, i, i
    BL     printf
    ADD    i, i, #1
    CMP    i, #10
    BLT    loop
    LDMFD  sp!, {i, pc}

print_string
    DCB    "Square of %d is %d\n", 0

    END
```

Instruction Scheduling

Instruction address	<i>pc</i>	<i>pc-4</i>	<i>pc-8</i>	<i>pc-12</i>	<i>pc-16</i>
Action	Fetch	Decode	ALU	LS1	LS2

ADD r0, r0, r1
 ADD r0, r0, r2

LDR r1, [r2, #4]
 ADD r0, r0, r1

Pipeline	Fetch	Decode	ALU	LS1	LS2
Cycle 1	...	ADD	LDR	...	
Cycle 2		...	ADD	LDR	...
Cycle 3		...	ADD	—	LDR

Instruction Scheduling (Cont.)

```
LDRB    r1, [r2, #1]
ADD     r0, r0, r2
EOR     r0, r0, r1
```

Pipeline
Cycle 1
Cycle 2
Cycle 3
Cycle 4

Fetch	Decode	ALU	LS1	LS2
EOR	ADD	LDRB	...	
...	EOR	ADD	LDRB	...
	...	EOR	ADD	LDRB
	...	EOR	—	ADD

```
MOV     r1, #1
B       case1
AND     r0, r0, r1
EOR     r2, r2, r3
...
case1   SUB     r0, r0, r1
```

Pipeline
Cycle 1
Cycle 2
Cycle 3
Cycle 4
Cycle 5

Fetch	Decode	ALU	LS1	LS2
AND	B	MOV	...	
EOR	AND	B	MOV	...
SUB	—	—	B	MOV
...	SUB	—	—	B
	...	SUB	—	—

Instruction Scheduling (Cont.)

- ARM cores have a pipeline architecture. The pipeline may delay the results of certain instructions for several cycles. If you use these results as source operands in a following instruction, the processor will insert stall cycles until the value is ready
- Load and multiply instructions have delayed results in many implementations
- You have two software methods available to remove interlocks following load instructions: You can preload so that loop i loads the data for loop $i + 1$, or you can unroll the loop and interleave the code for loops i and $i + 1$.

Register Allocation

- ARM has 14 available registers for general-purpose use: r0 to r12 and r14. The stack pointer r13 and program counter r15 cannot be used for general-purpose data. Operating system interrupts often assume that the user mode r13 points to a valid stack, so don't be tempted to reuse r13.
- If you need more than 14 local variables, swap the variables out to the stack, working outwards from the innermost loop

Register Allocation (Cont.)

- Use register names rather than physical register numbers when writing assembly routines. This makes it easier to reallocate registers and to maintain the code
- To ease register pressure you can sometimes store multiple values in the same register. For example, you can store a loop counter and a shift in one register. You can also store multiple pixels in one register

Conditional Execution

```
if (i<10)
{
    c = i + '0';
}
else
{
    c = i + 'A'-10;
}
```

```
CMP        i, #10
ADDLO      c, i, #'0'
ADDHS      c, i, #'A'-10
```

Conditional Execution

- You can implement most if statements with conditional execution. This is more efficient than using a conditional branch

Loop Constructs

```
    MOV i, N
loop
    ; loop body goes here and i=N,N-1,...,1
    SUBS i, i, #1
    BGT loop
```

```
    SUBS i, N, #1
loop
    ; loop body goes here and i=N-1,N-2,...,0
    SUBS i, i, #1
    BGE loop
```

Loop Constructs (Cont.)

```
    MOV i, N
loop
    ; loop body goes here and iterates (round up)(N/3) times
    SUBS i, i, #3
    BGT loop
```

Loop Constructs (Cont.)

- ARM requires two instructions to implement a counted loop: a subtract that sets flags and a conditional branch
- Unroll loops to improve loop performance. Do not overunroll because this will hurt cache performance. Unrolled loops may be inefficient for a small number of iterations. You can test for this case and only call the unrolled loop if the number of iterations is large

Loop Constructs (Cont.)

- Nested loops only require a single loop counter register, which can improve efficiency by freeing up registers for other uses
- ARM can implement negative and logarithmic indexed loops efficiently

Bit Manipulation

- The ARM can pack and unpack bits efficiently using logical operations and the barrel shifter
- To access bitstreams efficiently use a 32-bit register as a bitbuffer. Use a second register to keep track of the number of valid bits in the bitbuffer
- To decode bitstreams efficiently, use a lookup table to scan the next N bits of the bitstream. The lookup table can return codes of length at most N bits directly, or return an escape character for longer codes.

Efficient Switches

- Make sure the switch value is in the range $0 \leq x < N$ for some small N . To do this you may have to use a hashing function
- Use the switch value to index a table of function pointers or to branch to short sections of code at regular intervals. The second technique is position independent; the first isn't

Handling Unaligned Data

- If performance is not an issue, access unaligned data using multiple byte loads and stores. This approach accesses data of a given endianness regardless of the pointer alignment and the configured endianness of the memory system
- If performance is an issue, then use multiple routines, with a different routine optimized for each possible array alignment. You can use the assembler MACRO directive to generate these routines automatically

Thank you